

Please execute the code below and observe the output you get. Also, please learn

how to use each of these statements to get a similar task done.

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
%matplotlib inline
df =
pd.read_csv('https://raw.githubusercontent.com/jackiekazil/data-wrangling/master/data-text.csv')
df.head(2)
df1 =
pd.read_csv('https://raw.githubusercontent.com/kjam/data-wranglingpycon/master/data/berlin_weather_oldest.csv')
df1.head(2)
```

```
In [4]: #fetching data from online CSV files
import pandas as pd
df=pd.read_csv('https://raw.githubusercontent.com/jackiekazil/data-wrangling/master/data-text.csv')
df.head(2)
```

Out[4]:

	Indicator	PUBLISH STATES	Year	WHO region	World Bank income group	Country	Sex	Display Value	Numeric	Low	High	Col
0	Life expectancy at birth (years)	Published	1990	Europe	High-income	Andorra	Both sexes	77	77.0	NaN	NaN	
1	Life expectancy at birth (years)	Published	2000	Europe	High-income	Andorra	Both sexes	80	80.0	NaN	NaN	

```
In [3]: df1=pd.read_csv('https://raw.githubusercontent.com/kjam/data-wrangling-pycon/master/data/berlin_weather_oldest.csv')
df1.head(2) #fetching 2 rows from df
```

Out[3]:

	STATION	STATION_NAME	DATE	PRCP	SNWD	SNOW	TMAX	TMIN	WDFG	F
0	GHCND:GME00111445	BERLIN TEMPELHOF GM	19310101	46	-9999	-9999	-9999	-11	-9999	
1	GHCND:GME00111445	BERLIN TEMPELHOF GM	19310102	107	-9999	-9999	50	11	-9999	

2 rows × 21 columns

1. Get the Metadata from the above files.

```
In [5]: #fetching metadata from the 2 dataframes
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 4656 entries, 0 to 4655
Data columns (total 12 columns):
Indicator                4656 non-null object
PUBLISH STATES           4656 non-null object
Year                    4656 non-null int64
WHO region               4656 non-null object
World Bank income group  4656 non-null object
Country                 4656 non-null object
Sex                     4656 non-null object
Display Value            4656 non-null int64
Numeric                 4656 non-null float64
Low                      0 non-null float64
High                     0 non-null float64
Comments                 0 non-null float64
dtypes: float64(4), int64(2), object(6)
memory usage: 436.6+ KB
```

```
In [6]: df1.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 117208 entries, 0 to 117207
Data columns (total 21 columns):
STATION                117208 non-null object
STATION_NAME           117208 non-null object
DATE                  117208 non-null int64
PRCP                  117208 non-null int64
SNWD                  117208 non-null int64
SNOW                  117208 non-null int64
TMAX                  117208 non-null int64
TMIN                  117208 non-null int64
WDFG                  117208 non-null int64
PGTM                  117208 non-null int64
WSFG                  117208 non-null int64
WT09                  117208 non-null int64
WT07                  117208 non-null int64
WT01                  117208 non-null int64
WT06                  117208 non-null int64
WT05                  117208 non-null int64
WT04                  117208 non-null int64
WT16                  117208 non-null int64
WT08                  117208 non-null int64
WT18                  117208 non-null int64
WT03                  117208 non-null int64
dtypes: int64(19), object(2)
memory usage: 18.8+ MB
```

2. Get the row names from the above files.

```
In [7]: #Getting the row names from the 2 dataframes
df.index.values
```

```
Out[7]: array([ 0, 1, 2, ..., 4653, 4654, 4655], dtype=int64)
```

```
In [8]: df1.index.values
```

```
Out[8]: array([ 0, 1, 2, ..., 117205, 117206, 117207], dtype=int64)
```

3. Change the column name from any of the above file.

```
In [11]: # Renaming Indicator to indicator_id
df.rename(columns={'Indicator': 'indicator_id'})
```

Out[11]:

	indicator_id	PUBLISH STATES	Year	WHO region	World Bank income group	Country	Sex	Display Value	Numeric	Low	High	C
0	Life expectancy at birth (years)	Published	1990	Europe	High-income	Andorra	Both sexes	77	77.0	NaN	NaN	
1	Life expectancy at birth (years)	Published	2000	Europe	High-income	Andorra	Both sexes	80	80.0	NaN	NaN	

4. Change the column name from any of the above file and store the changes made permanently.

```
In [12]: # Renaming Indicator to indicator_id and making it permanent
df.rename(columns={'Indicator': 'indicator_id'}, inplace=True)
df.head(2)
```

Out[12]:

	indicator_id	PUBLISH STATES	Year	WHO region	World Bank income group	Country	Sex	Display Value	Numeric	Low	High	C
0	Life expectancy at birth (years)	Published	1990	Europe	High-income	Andorra	Both sexes	77	77.0	NaN	NaN	
1	Life expectancy at birth (years)	Published	2000	Europe	High-income	Andorra	Both sexes	80	80.0	NaN	NaN	

5. Change the names of multiple columns.

```
In [13]: # Changing column names ('PUBLISH STATES', 'WHO region') to ('Publication Status', 'WHO Region')
df.rename(columns={'PUBLISH STATES': 'Publication Status', 'WHO region': 'WHO Region'})
df.head(2)
```

Out[13]:

	indicator_id	Publication Status	Year	WHO Region	World Bank income group	Country	Sex	Display Value	Numeric	Low	High
0	Life expectancy at birth (years)	Published	1990	Europe	High-income	Andorra	Both sexes	77	77.0	NaN	NaN
1	Life expectancy at birth (years)	Published	2000	Europe	High-income	Andorra	Both sexes	80	80.0	NaN	NaN

6. Arrange values of a particular column in ascending order

```
In [16]: # Arranging columns based on Year column
df.sort_values('Year').head()
```

Out[16]:

	indicator_id	Publication Status	Year	WHO Region	World Bank income group	Country	Sex	Display Value	Numeric	Low	Hi
0	Life expectancy at birth (years)	Published	1990	Europe	High-income	Andorra	Both sexes	77	77.0	NaN	N
1270	Life expectancy at birth (years)	Published	1990	Europe	High-income	Germany	Male	72	72.0	NaN	N
3193	Life expectancy at birth (years)	Published	1990	Europe	Lower-middle-income	Republic of Moldova	Male	65	65.0	NaN	N
3194	Life expectancy at birth (years)	Published	1990	Europe	Lower-middle-income	Republic of Moldova	Both sexes	68	68.0	NaN	N
3197	Life expectancy at age 60 (years)	Published	1990	Europe	Lower-middle-income	Republic of Moldova	Male	15	15.0	NaN	N

7. Arrange multiple column values in ascending order.

In [25]: *# Unable to get the same results because there is no combination available to get*
`df.sort_values(['Publication Status', 'indicator_id', 'Country', 'Year'], ascending=[`

Out[25]:

	Country	indicator_id	Publication Status	Year	WHO Region	World Bank income group	Sex	Display Value	Numer
554	Afghanistan	Life expectancy at birth (years)	Published	1990	Eastern Mediterranean	Low-income	Both sexes	49	49
965	Afghanistan	Life expectancy at birth (years)	Published	1990	Eastern Mediterranean	Low-income	Male	49	49
1792	Afghanistan	Life expectancy at birth (years)	Published	1990	Eastern Mediterranean	Low-income	Female	50	50
146	Afghanistan	Life expectancy at birth	Published	2000	Eastern Mediterranean	Low-income	Both sexes	55	55

8. Make countryas the first column of the dataframe.

In [22]: *# Taking the column into a List, and repositioning the column by making changes*
`cols=df.columns.tolist()
cols.remove("Country")
cols = ["Country"] + cols
df=df[cols]
df.head(5)`

Out[22]:

	Country	indicator_id	Publication Status	Year	WHO Region	World Bank income group	Sex	Display Value	Numeric	Low
0	Andorra	Life expectancy at birth (years)	Published	1990	Europe	High-income	Both sexes	77	77.0	NaN
1	Andorra	Life expectancy at birth (years)	Published	2000	Europe	High-income	Both sexes	80	80.0	NaN
2	Andorra	Life expectancy at age 60 (years)	Published	2012	Europe	High-income	Female	28	28.0	NaN
3	Andorra	Life expectancy at age 60 (years)	Published	2000	Europe	High-income	Both sexes	23	23.0	NaN
4	United Arab Emirates	Life expectancy at birth (years)	Published	2012	Eastern Mediterranean	High-income	Female	78	78.0	NaN

9. Get the column array using a variable

```
In [23]: #fetching column name into variable
who_region=df["WHO Region"].as_matrix()
who_region
```

C:\Users\manoj\Anaconda3\lib\site-packages\ipykernel_launcher.py:1: FutureWarning: Method .as_matrix will be removed in a future version. Use .values instead.
 """Entry point for launching an IPython kernel.

```
Out[23]: array(['Europe', 'Europe', 'Europe', ..., 'Africa', 'Africa', 'Africa'],
              dtype=object)
```

10. Get the subset rows 11, 24, 37

```
In [26]: df.iloc[[11,24,37]]
```

```
Out[26]:
```

	Country	indicator_id	Publication Status	Year	WHO Region	World Bank income group	Sex	Display Value	Numeric	Low
11	Austria	Life expectancy at birth (years)	Published	2012	Europe	High-income	Female	83	83.0	NaN
24	Brunei Darussalam	Life expectancy at age 60 (years)	Published	2012	Western Pacific	High-income	Female	21	21.0	NaN
37	Cyprus	Life expectancy at age 60 (years)	Published	2012	Europe	High-income	Female	26	26.0	NaN

11. Get the subset rows excluding 5, 12, 23, and 56

```
In [35]: # using the ~ and isin function to exclude the specific rows
df[~df.index.isin([5,12,23,56])]
```

```
Out[35]:
```

	Country	indicator_id	Publication Status	Year	WHO Region	World Bank income group	Sex	Display Value	Numer
0	Andorra	Life expectancy at birth (years)	Published	1990	Europe	High-income	Both sexes	77	77
1	Andorra	Life expectancy at birth (years)	Published	2000	Europe	High-income	Both sexes	80	80
2	Andorra	Life expectancy at age 60 (years)	Published	2012	Europe	High-income	Female	28	28
3	Andorra	Life expectancy at age 60	Published	2000	Europe	High-income	Both sexes	23	23

Load datasets from CSV

```

users=pd.read_csv('https://raw.githubusercontent.com/ben519/DataWrangling/master/Data/users.csv')
sessions=pd.read_csv('https://raw.githubusercontent.com/ben519/DataWrangling/master/Data/sessions.csv')
products=pd.read_csv('https://raw.githubusercontent.com/ben519/DataWrangling/master/Data/products.csv')
transactions=pd.read_csv('https://raw.githubusercontent.com/ben519/DataWrangling/master/Data/transactions.csv')
users.head()
sessions.head()
transactions.head()

```

In [80]: *#fetching CSV files from URLs into dataframes and getting first 2 rows*

```

users=pd.read_csv('https://raw.githubusercontent.com/ben519/DataWrangling/master/Data/users.csv')
sessions=pd.read_csv('https://raw.githubusercontent.com/ben519/DataWrangling/master/Data/sessions.csv')
products=pd.read_csv('https://raw.githubusercontent.com/ben519/DataWrangling/master/Data/products.csv')
transactions=pd.read_csv('https://raw.githubusercontent.com/ben519/DataWrangling/master/Data/transactions.csv')
users.head(2)

```

Out[80]:

	UserID	User	Gender	Registered	Cancelled
0	1	Charles	male	2012-12-21	NaN
1	2	Pedro	male	2010-08-01	2010-08-08

In [37]: sessions.head(2)

Out[37]:

	SessionID	SessionDate	UserID
0	1	2010-01-05	2
1	2	2010-08-01	2

In [38]: products.head(2)

Out[38]:

	ProductID	Product	Price
0	1	A	14.16
1	2	B	33.04

In [39]: transactions.head(2)

Out[39]:

	TransactionID	TransactionDate	UserID	ProductID	Quantity
0	1	2010-08-21	7.0	2	1
1	2	2011-05-26	3.0	4	1

12. Join users to transactions, keeping all rows from transactions and only matching rows from users (left join)

In [42]: *# Using merge with parameter Left to get all rows from transactions and only match*
`transactions.merge(users, how='left', on='UserID')`

Out[42]:

	TransactionID	TransactionDate	UserID	ProductID	Quantity	User	Gender	Registered	Canceled
0	1	2010-08-21	7.0	2	1	NaN	NaN	NaN	
1	2	2011-05-26	3.0	4	1	Caroline	female	2012-10-23	201
2	3	2011-06-16	3.0	3	1	Caroline	female	2012-10-23	201
3	4	2012-08-26	1.0	2	3	Charles	male	2012-12-21	
4	5	2013-06-06	2.0	4	1	Pedro	male	2010-08-01	201
5	6	2013-12-23	2.0	5	6	Pedro	male	2010-08-01	201
6	7	2013-12-30	3.0	4	1	Caroline	female	2012-10-23	201
7	8	2014-04-24	NaN	2	3	NaN	NaN	NaN	
8	9	2015-04-24	7.0	4	3	NaN	NaN	NaN	
9	10	2016-05-08	3.0	4	4	Caroline	female	2012-10-23	201

13. Which transactions have a UserID not in users?

In [46]: *# same way as question 11, using ~ and isin function to exclude the rows in transactions*
`transactions[~transactions['UserID'].isin(users['UserID'])]`

Out[46]:

	TransactionID	TransactionDate	UserID	ProductID	Quantity
0	1	2010-08-21	7.0	2	1
7	8	2014-04-24	NaN	2	3
8	9	2015-04-24	7.0	4	3

14. Join users to transactions, keeping only rows from transactions and users that match via UserID (inner join)

In [47]: *# Using merge with parameter inner to get only matching rows from both the dataframes*
 transactions.merge(users, how='inner', on='UserID')

Out[47]:

	TransactionID	TransactionDate	UserID	ProductID	Quantity	User	Gender	Registered	Canceled
0	2	2011-05-26	3.0	4	1	Caroline	female	2012-10-23	201
1	3	2011-06-16	3.0	3	1	Caroline	female	2012-10-23	201
2	7	2013-12-30	3.0	4	1	Caroline	female	2012-10-23	201
3	10	2016-05-08	3.0	4	4	Caroline	female	2012-10-23	201
4	4	2012-08-26	1.0	2	3	Charles	male	2012-12-21	
5	5	2013-06-06	2.0	4	1	Pedro	male	2010-08-01	201
6	6	2013-12-23	2.0	5	6	Pedro	male	2010-08-01	201

15. Join users to transactions, displaying all matching rows AND all non-matching rows (full outer join)

In [48]: *# Using merge with parameter outer to get all rows from transactions and users*
 transactions.merge(users, how='outer', on='UserID')

Out[48]:

	TransactionID	TransactionDate	UserID	ProductID	Quantity	User	Gender	Registered	Canceled
0	1.0	2010-08-21	7.0	2.0	1.0	NaN	NaN	NaN	
1	9.0	2015-04-24	7.0	4.0	3.0	NaN	NaN	NaN	
2	2.0	2011-05-26	3.0	4.0	1.0	Caroline	female	2012-10-23	2
3	3.0	2011-06-16	3.0	3.0	1.0	Caroline	female	2012-10-23	2
4	7.0	2013-12-30	3.0	4.0	1.0	Caroline	female	2012-10-23	2
5	10.0	2016-05-08	3.0	4.0	4.0	Caroline	female	2012-10-23	2
6	4.0	2012-08-26	1.0	2.0	3.0	Charles	male	2012-12-21	
7	5.0	2013-06-06	2.0	4.0	1.0	Pedro	male	2010-08-01	2
8	6.0	2013-12-23	2.0	5.0	6.0	Pedro	male	2010-08-01	2
9	8.0	2014-04-24	NaN	2.0	3.0	NaN	NaN	NaN	
10	NaN	NaN	4.0	NaN	NaN	Brielle	female	2013-07-17	
11	NaN	NaN	5.0	NaN	NaN	Benjamin	male	2010-11-25	

16. Determine which sessions occurred on the same day each user registered

In [49]: *# Using merge with parameter innter to get all rows from transactions and users w*
`pd.merge(left=users, right=sessions, how='inner', left_on=['UserID', 'Registered'`

Out[49]:

	UserID	User	Gender	Registered	Cancelled	SessionID	SessionDate
0	2	Pedro	male	2010-08-01	2010-08-08	2	2010-08-01
1	4	Brielle	female	2013-07-17	NaN	9	2013-07-17

17. Build a dataset with every possible (UserID, ProductID) pair (cross join)

```
In [83]: # Fetching the required columns into 2 dataframes
df_users=pd.DataFrame(users['UserID'],columns=['UserID'])
df_products=pd.DataFrame(products['ProductID'],columns=["ProductID"])

# adding a row called key with value as 1
df_users['key']=1
df_products['key']=1

# Using the merge with default how parameter to get the required output
df_U_P = pd.merge(df_users,df_products,on='key')[["UserID","ProductID"]]
df_U_P
```

```
Out[83]:
```

	UserID	ProductID
0	1	1
1	1	2
2	1	3
3	1	4
4	1	5
5	2	1
6	2	2
7	2	3
8	2	4
9	2	5
10	3	1
11	3	2
12	3	3
13	3	4
14	3	5
15	4	1
16	4	2
17	4	3
18	4	4
19	4	5
20	5	1
21	5	2
22	5	3
23	5	4
24	5	5

18. Determine how much quantity of each product was purchased by each user

```
In [88]: # Using merge, groupby and sum to get the quantity of each product was purchased
pd.merge(df_U_P, transactions, how='left').groupby(["UserID", "ProductID"])["Quantity"]
```

```
Out[88]: UserID  ProductID
1           1           0.0
           2           3.0
           3           0.0
           4           0.0
           5           0.0
2           1           0.0
           2           0.0
           3           0.0
           4           1.0
           5           6.0
3           1           0.0
           2           0.0
           3           1.0
           4           6.0
           5           0.0
4           1           0.0
           2           0.0
           3           0.0
           4           0.0
           5           0.0
5           1           0.0
           2           0.0
           3           0.0
           4           0.0
           5           0.0
Name: Quantity, dtype: float64
```

**19. For each user, get each possible pair of pair transactions
(TransactionID1, TransacationID2)**

In [89]: *# Using the merge on transactions twice based on "UserID" to get possible pair of*
 pd.merge(transactions, transactions, on='UserID')

Out[89]:

	TransactionID_x	TransactionDate_x	UserID	ProductID_x	Quantity_x	TransactionID_y	Transact
0	1	2010-08-21	7.0	2	1	1	2
1	1	2010-08-21	7.0	2	1	9	2
2	9	2015-04-24	7.0	4	3	1	2
3	9	2015-04-24	7.0	4	3	9	2
4	2	2011-05-26	3.0	4	1	2	2
5	2	2011-05-26	3.0	4	1	3	2
6	2	2011-05-26	3.0	4	1	7	2
7	2	2011-05-26	3.0	4	1	10	2
8	3	2011-06-16	3.0	3	1	2	2
9	3	2011-06-16	3.0	3	1	3	2
10	3	2011-06-16	3.0	3	1	7	2
11	3	2011-06-16	3.0	3	1	10	2
12	7	2013-12-30	3.0	4	1	2	2
13	7	2013-12-30	3.0	4	1	3	2
14	7	2013-12-30	3.0	4	1	7	2
15	7	2013-12-30	3.0	4	1	10	2
16	10	2016-05-08	3.0	4	4	2	2
17	10	2016-05-08	3.0	4	4	3	2
18	10	2016-05-08	3.0	4	4	7	2
19	10	2016-05-08	3.0	4	4	10	2
20	4	2012-08-26	1.0	2	3	4	2
21	5	2013-06-06	2.0	4	1	5	2
22	5	2013-06-06	2.0	4	1	6	2
23	6	2013-12-23	2.0	5	6	5	2
24	6	2013-12-23	2.0	5	6	6	2
25	8	2014-04-24	NaN	2	3	8	2

20. Join each user to his/her first occuring transaction in the transactions table

In [91]: *# Using merge, left and first method to get each user to his/her first occurring transaction*
 pd.merge(users, transactions.groupby('UserID').first(), how='left', on='UserID')

Out[91]:

	UserID	User	Gender	Registered	Cancelled	TransactionID	TransactionDate	ProductID	Quantity
0	1	Charles	male	2012-12-21	NaN	4.0	2012-08-26	2.0	
1	2	Pedro	male	2010-08-01	2010-08-08	5.0	2013-06-06	4.0	
2	3	Caroline	female	2012-10-23	2016-06-07	2.0	2011-05-26	4.0	
3	4	Brielle	female	2013-07-17	NaN	NaN	NaN	NaN	
4	5	Benjamin	male	2010-11-25	NaN	NaN	NaN	NaN	

This is for the bonus points

```
In [111]: def removeYears(item):
            if type(item) == str:
                if '(years)' in item:
                    return item.replace("(years)", "")
                else:
                    return item
            else:
                return item
df1 = df.applymap(removeYears)
df1
```

6	Antigua and Barbuda	Life expectancy at age 60	Published	1990	Americas	High-income	Male	17	17
7	Antigua and Barbuda	Life expectancy at age 60	Published	2012	Americas	High-income	Both sexes	22	22
8	Australia	Life expectancy at birth	Published	2012	Western Pacific	High-income	Male	81	81
9	Australia	Life expectancy at birth	Published	2000	Western Pacific	High-income	Both sexes	80	80

In []: